

CS553 Homework #3

Understanding the Cost of Cloud Computing

Instructions:

- *Assigned date: Saturday September 13th, 2025*
- *Due date: 11:59PM on Thursday September 25th, 2025*
- *Maximum Points: 100*
- *This assignment can be done in groups of 2 students*
- *Please post questions on Canvas*
- *Only a softcopy submission is required through GIT: push changes to GIT repository, and email confirmation will be sent to your HAWK email address at the deadline*
- *Late submission will be penalized at 20%; late submissions beyond 24 hours will receive 0 points*

1. Introduction

You are hired by a startup company who is considering to use cloud computing instead of building its own infrastructure. There is consensus that a cloud computing software stack at the layer of IaaS will be used, but it's not clear whether the computing resources should be rented from a public cloud on-demand, or whether a private cloud should be purchased. You are tasked to find the cost breakdown of a private cloud, and compare that to what Amazon would charge. You can find many instance types defined at <http://aws.amazon.com/ec2/instance-types/>, and their prices are set at <http://aws.amazon.com/ec2/pricing/>. For pricing purposes, please stick to Linux on-demand pricing. There are a variety of Amazon calculators such as <https://calculator.aws/#/> and <https://calculator.aws/#/createCalculator/S3>, please use them if you find them useful.

Since you have to estimate the cost of the hardware when building a private cloud, you can use hardware prices found at ThinkMate website (<https://www.thinkmate.com>) as good sources for server hardware (for configuration #1 and #3). For configuration #2, you will need to use the Apple website (<https://www.apple.com/mac-mini/>). You must include a printout of your shopping cart in your final writeup report for this assignment; include this as an appendix at the end of your report.

You are to estimate the cost of different configurations for 3 different set of requirements; compute prices for a 5-year period:

- Configuration 1: Hadoop/Spark Cluster with 1M-cores, 2048TB memory, 59.375PB HDD, and 50Gb/s Ethernet Fat-Tree network (each VM should be equivalent to the c6id.metal instance); in addition to the compute resources, a 100PB distributed storage shared across the entire cloud should be procured, with the expectation that 100PB of data will be read and written to S3 every year from outside of Amazon with enough capacity for 1GB/sec throughput (for pricing comparison, see S3 Standard). For EC2, you can use the lowest advertised price for this instance with up to a 5 year term commitment, if such pricing is available.
- Configuration 2: Support 1K application developers who are designing MacOS and iPad OS applications. They require a MacOS system with 14-cores, 48GB RAM, 1TB storage, and 10Gb/s network (Amazon has mac-m4pro.metal instances that have everything you need except the 1TB storage, which you can provision through EBS). The developers work 40 hours/week, 48 weeks/year (they get 4 weeks of vacation per year). You must use on-demand EC2 pricing as developers are expected to provision their systems at the beginning of each working day, and release their systems at the end of each working day.
- Configuration 3: Monero (XMR) crypto currency mining; you have an investor who has \$10M to buy

hardware to mine Monero Coin XMR (and pay for maintenance / sys admin, power, and cooling), or rent resources from Amazon EC2 to mine XMR Coin. Configure the best hardware you can from ThinkMate. For buying hardware solution, make sure to leave funds to pay for power, cooling, and system administrator. XMR Coin mining can be done on any compute hardware (CPUs or GPUs), but you will likely find that it's most profitable to mine using CPUs. Since XMR mining is compute intensive, your processor will likely require to have as many cores as possible; if GPU mining is used, you might be able to lower the processing and memory requirements to a minimum on the host machine. Identify the best Amazon instance (you must use Spot Instances to make sure you get the best hardware for the cheapest price); although spot pricing fluctuates over time, you can assume the spot price will remain fixed for the duration of your evaluation. For the purchase of the hardware scenario, you are free to locate the hardware in any state in the USA (for a full list of average electricity cost by state, see <https://www.chooseenergy.com/electricity-rates-by-state/>); since this will be a business venture, use the business electricity rates. If electricity is too expensive to make a profit, invest part of the \$1M in solar power (solar panels), and estimate the amount of energy you can extract. For an overview of various CPUs and their respective hashrates (the higher the hashrates, the more XMR Coin that can be mined), see <https://hashrate.no/coins/XMR/benchmarks>. Once you have a hashrate, you can estimate how much money can be made mining XMR by using an online calculator such as <https://minerstat.com/coin/XMR>. The mining calculator gives an instantaneous mining number, although in reality the amount of coin that can be mined would vary based on many factors (hash rate, hash difficulty, fees, etc). The profit similarly can vary based on the XMR Coin pricing, which can vary wildly. When computing the mining coins and expected profit, you can use the calculator above to compute it for a 5-year period, assuming the mining continues at the same rate, and the price remains at the same level. Your task is to compute the amount of profit that is expected after \$10M is invested in buying hardware and running it for 5-years, vs. renting the hardware from Amazon. It's possible that the profits you make will be less than the original investment (especially with the Amazon scenario).

2. What you will submit?

Your deliverables for this project are to be written in a report, which will include the following:

- Report: A written document (typed, named hw3-report.pdf) describing your answers to the above questions.
- Compare the costs of the 3 different configurations between the public cloud (Amazon AWS) and the private cloud
 - you may assume a 5-year amortization cost
 - you will have to factor in things other than hardware, such as cooling, power, administration costs, network infrastructure (e.g. switches); you can assume 1 system administrator is needed for every 500 servers
 - show your data in three different tables with the costs of each of the 3 configurations, broken down by components (e.g. servers, network switches, cables, racks, cooling, power, administration, etc)
 - summarize your data in a 4th table, comparing the public cloud cost to the private cloud cost
- Explain in words if it is better to rent or buy, and by how much.
- Include your shopping cart of the 3 configurations. Your submission should be a single large PDF file, starting with your report, and followed by the shopping carts.

Table 1: Table summarizing each of the three configurations; please include 1 table per configuration, for a total of 3 tables

	Description	Price per Item	Quantity	Total Price
Compute Servers				
Network Switches				
Network Cables				
Racks				
Storage Servers				
Electric Power				
Cooling				
Administration				
TOTAL	N/A	N/A	N/A	

Table 2: Summary table comparing the 2 configurations between the public and private cloud; your cost of power, cooling, and administration should be to cover 5 years of costs

	Configuration 1	Configuration 2
Public Cloud (including EC2 and S3) Cost over 5 years		
Private Cloud cost over 5 years		

Table 3: Summary table comparing the profits with configuration #3 between the public and private cloud over a 5-year period

	Configuration 3
Public Cloud Mining Profit over 5 years	
Private Cloud Mining Profit over 5 years	

You will have to submit your solution to a private git repository created for you at <https://classroom.github.com/a/ne-aZh02>. You will have to firstly clone the repository. Then you will have to add or update your source code, documentation and report. Your solution will be collected automatically after the deadline. If you want to submit your homework later, you will have to push your final version to your GIT repository and you will have let the TA know of it through email. There is no need to submit anything on BB for this assignment. If you cannot access your repository contact the TAs. You can find a git cheat sheet here: <https://www.git-tower.com/blog/git-cheat-sheet/>

Submit step: code/report through GIT.

Grades for late programs will be lowered 20%; no submission will be accepted more than 24 hours late